

Lab Report in Bayesian Learning

Laboration 3

TDDE07

David Grönberg, Erik Lindgren
Davgr686, Erili040



17-05-2020

Contents

1	Assignments	1
1.1	Assignment 1	1
1.1.1	a)	1
1.1.2	b)	2
1.1.3	c)	2
1.2	Assignment 2	4
1.2.1	a)	4
1.2.2	b)	4
1.2.3	c)	6
1.2.4	d)	8
2	Code Appendix	10
2.1	Assignment 1	10
2.1.1	a	10
2.1.2	b	11
2.1.3	c	11
2.2	Assignment 2	11

List of Figures

1.1	Trajectories of the sampled Markov chains for μ (left) and σ^2 (right).	1
1.2	Fitted density of the Mixture of normal models(red) plotted over the precipitation data(black).	2
1.3	Graphical comparison between an normal density(blue) and Mixture normal density(red) on precipitation kernel density estimate(black)	3
1.4	Barplot showing the β values for each covariate	4
1.5	Histogram of β draws	6
1.6	Draws from the posterior of β	7
1.7	Cumulative mean of the draws from the posterior of β	8
1.8	Predictive distribution of $nBids$	9

List of Tables

1. Assignments

1.1 Assignment 1

We are given a dataset containing daily measurements of precipitation between the years of 1948 to 1983. In this assignment we will analyse the data using first a normal model and then a mixture normal model. The results of the models will then be compared graphically.

1.1.1 a)

We assume that the daily precipitation are independent normally distributed, where $\mu \sim N(\mu_0, \tau_0^2)$ and $\sigma^2 \sim Inv - \chi^2(v_0, \sigma_0^2)$. The first task of the assignment is to implement a Gibbs sampler that simulates from the joint posterior $p(\mu_0, \sigma_0^2 | y_1, \dots, y_n)$.

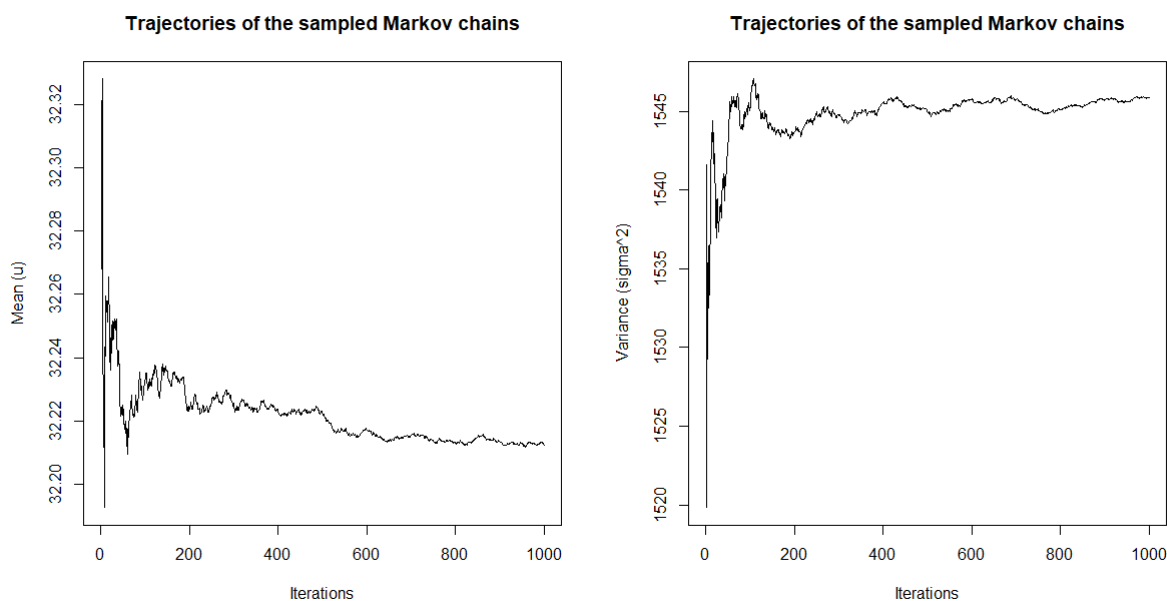


Figure 1.1: Trajectories of the sampled Markov chains for μ (left) and σ^2 (right).

In figure 1.1 we can see the trajectories of the sampled Markov chains for μ and σ^2 . From the figure we can see that both of the trajectories converge, μ towards 32 and σ^2 towards 1546.

1.1.2 b)

In the next task we instead assume that the daily precipitation follow an iid two-component mixture of normals model:

$$p(y_i|\mu, \sigma^2, \pi) = \pi N(y_i|\mu_1, \sigma_1^2) + (1 - \pi)N(y_i|\mu_2, \sigma_2^2)$$

where $\mu = (\mu_1, \mu_2)$ and $\sigma^2 = (\sigma_1^2, \sigma_2^2)$.




Figure 1.2: Fitted density of the Mixture of normal models(red) plotted over the precipitation data(black).

In figure 1.2 we can see that the mixture model fits nicely to the data. The model is created through 100 Gibbs sampling draws and with two mixture components. The prior mean and std of mu is set to 0 and 10. The Dirichlet(alpha) is set to 10 and prior sigma is set to the variance of the input data with 4 degrees of freedom.

1.1.3 c)

The last task in assignment 1 is to make a graphical comparison between the Normal model and the Mixture normal model. From figure 1.3 we can see that the fitted density of the Mixture normal model fits the data significantly better than the normal model.

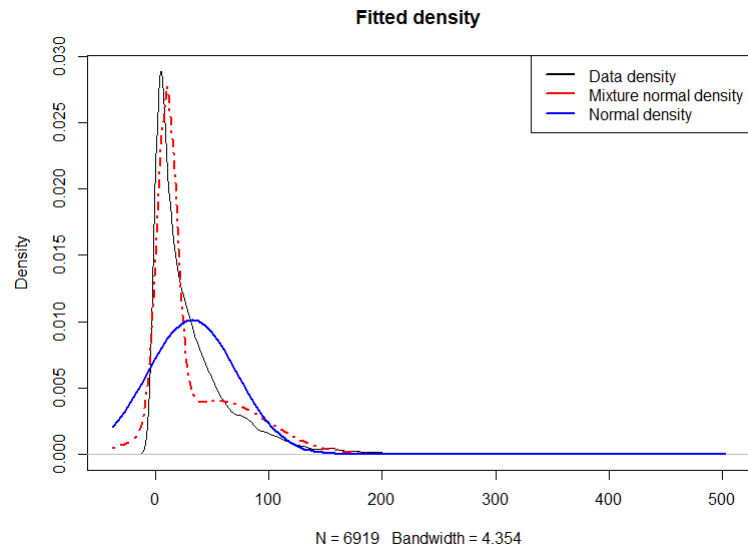


Figure 1.3: Graphical comparison between an normal density(blue) and Mixture normal density(red) on precipitation kernel density estimate(black)

1.2 Assignment 2

The dataset *eBayNumberOfBidderData.dat* contains observations from 1000 eBay auctions of coins. The response variable *nBids* denotes the number of bids in each auction. *nBids* can be modeled with the poisson regression model:

$$y_i|\beta \sim \text{Poisson}[\exp(x_i^T \beta)], i = 1, \dots, n$$

where y_i corresponds to the number of bids for the i th observation and x_i are the covariates for the i th observation.

1.2.1 a)

To obtain the maximum likelihood estimator of β in the Poisson regression model, we used the built in *glm()* function.

Figure 1.4 shows a barplot with the covariates beta values. As the plot shows, The covariates *Intercept*, *MinBidShare*, *Sealed*, *VerifyID* and *MajBlem* are significant in predicting the number of bids.




Figure 1.4: Barplot showing the β values for each covariate

1.2.2 b)

We conduct Bayesian analysis of the Poisson regression. We use the Zellners g-prior:

$$\beta \sim N[0, 100 * (X^T X)^{-1}]$$

where X is the $n \times p$ covariate matrix.

We assume that the posterior density is approximately multivariate normal

$$\beta|y \sim N(\hat{\beta}, J_y^{-1}(\hat{\beta}))$$

where $\hat{\beta}$ is the posterior mode and $J_y(\hat{\beta})$ is the negative Hessian at the posterior mode. These are obtained by numerical optimization using the *optim()* function in R. The log likelihood for the Poisson regression is expressed as:

$$\log L(\theta \mid X, Y) = \sum_{i=1}^m \left(y_i \theta' x_i - e^{\theta' x_i} - \log(y_i!) \right)$$

and because $\log(y_i!)$ does not contain θ , we may drop it from the expression. Leaving us with the following expression:

$$\ell(\theta \mid X, Y) = \sum_{i=1}^m \left(y_i \theta' x_i - e^{\theta' x_i} \right)$$

A function that calculates the posterior using the log likelihood and g-prior is used in the *optim()* function. Figure 1.5 shows the histogram of the β values when drawing from the multivariate normal model after obtaining values for $\hat{\beta}$ and $J_y(\hat{\beta})$.

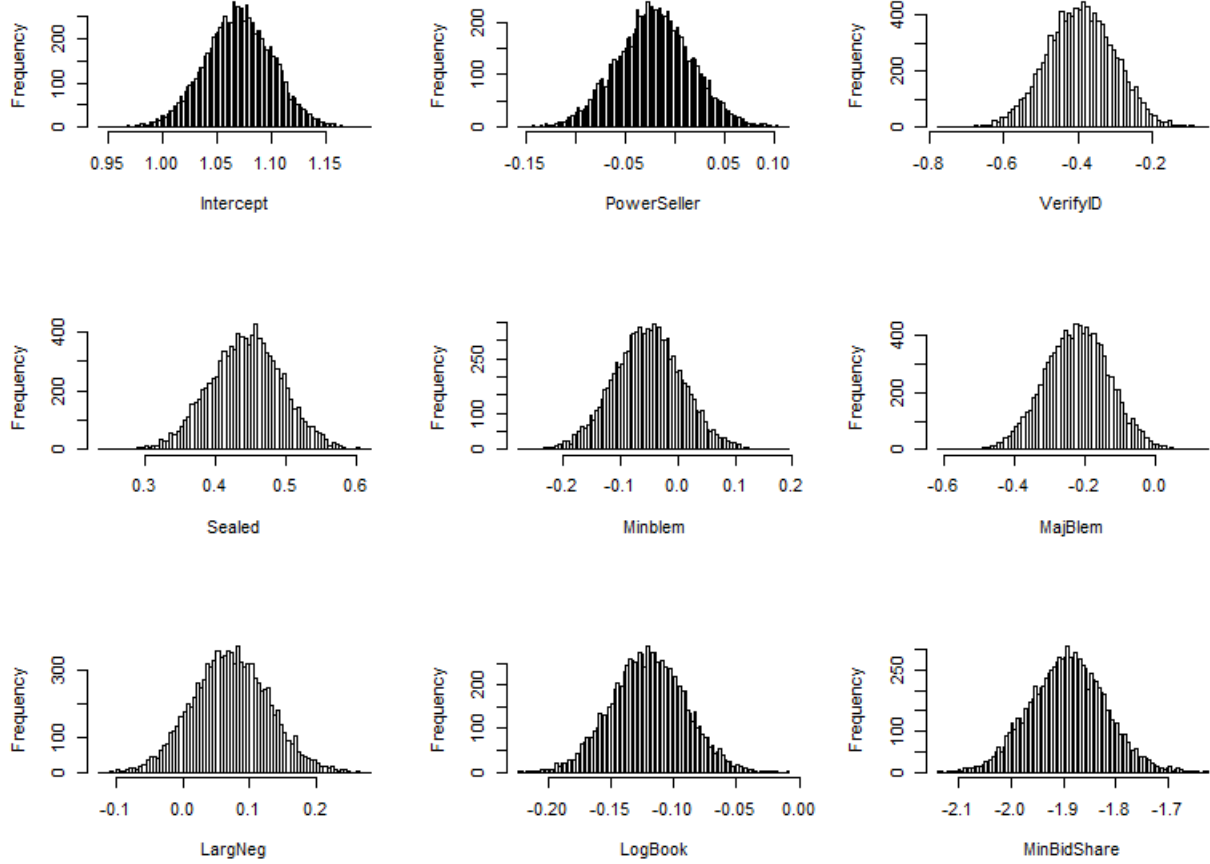


Figure 1.5: Histogram of β draws

1.2.3 c)

Next, we create a general function that takes an arbitrary posterior density and uses the Metropolis algorithm to generate random draws from the given posterior density. We let the proposal density be the multivariate normal density:

$$\theta_p | \theta^{(i-1)} \sim N(\theta^{(i-1)}, c*)$$

where c is a tuning parameter and $c* = J_y^{-1}(\hat{\beta})$. We use the created Metropolis function and input the log posterior of the Poisson regression to sample from the posterior of β

Figure 1.6 shows 10,000 draws of the β posterior using the Metropolis function. The draws reaches a stationary distribution after about 200 draws. Figure 1.7 shows the cumulative mean of the β posterior draws (represented by the black curve), and the maximum likelihood estimate of each β , obtained by the *glm()* model. The cumulative mean converges to the MLE, however, because of the "burn-in", it takes a longer time, compared to calculating the cumulative mean from sample 200-10,000.

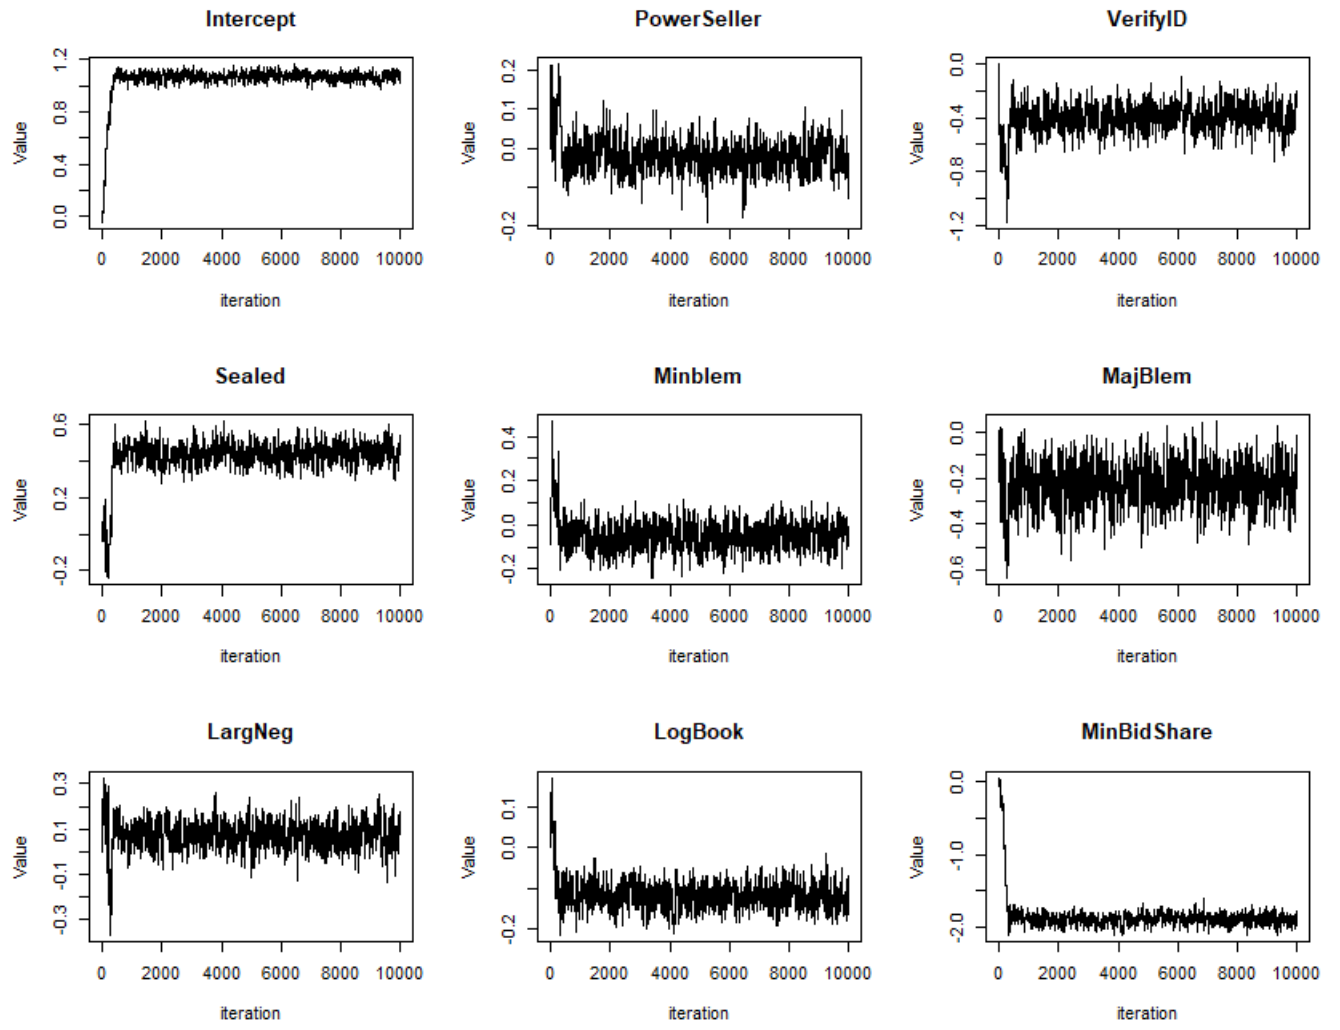


Figure 1.6: Draws from the posterior of β

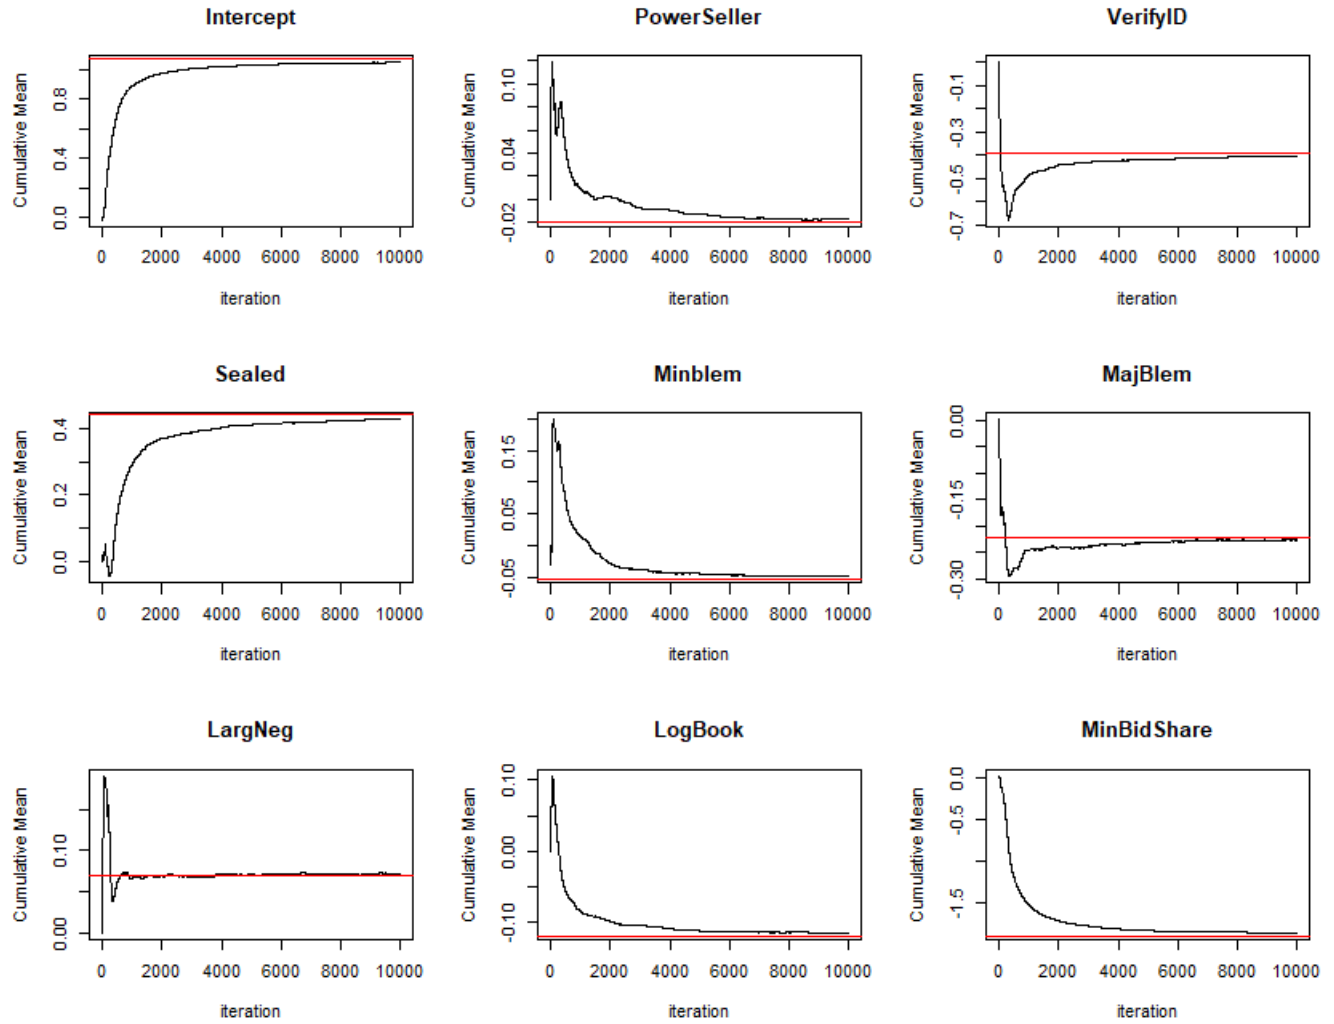


Figure 1.7: Cumulative mean of the draws from the posterior of β

1.2.4 d)

Next, we use the draws from c) to simulate from the predictive distribution of the number of bidders with the following characteristics:

Powerseller = 1

VerifyID = 1

Sealed = 1

MinBlem = 0

MajBlem = 0

LargeNeg = 0

LogBook = 1

MinBidShare = 0.5

Figure 1.8 shows a barplot of the predictive distribution of the number of bidders when a auction has the characteristics described above. The cumulative mean obtained from c) is used as the β . The probability of having zero bidders with the described characteristics is around 30%.




Figure 1.8: Predictive distribution of $nBids$

2. Code Appendix

2.1 Assignment 1

2.1.1 a

```
1 library(mvtnorm)
2 df = read.table("rainfall.dat", header = TRUE)
3 n = length(df$X136)
4
5 #prior
6 tau_0 = 1
7 u_0 = 32
8 v_0 = 1
9 sigma_0 = 1
10 sigma = 1
11
12 x = mean(df$X136)
13 v_n = v_0 + n
14 draws = matrix(0,1000,2)
15
16 for (j in 1:1000){
17   tau_n = 1/((n/sigma)+(1/tau_0))
18   w = (n/sigma)*tau_n
19   u_n = w*x + (1-w)*u_0
20
21   u = rnorm(1,u_n,tau_n)
22   sigma = (v_n * (v_0*sigma_0 + sum((df$X136 - u )^2))/(n*v_0)) / rchisq(1,df=v_n)
23
24   draws[j,1] = u
25   draws[j,2] = sigma
26 }
27 hist(draws[,1],50) #u
28 hist(draws[,2],50) #sigma^2
29
30 plot(draws[,1],type='l')
31 plot(draws[,2],type='l')
32
33 TrajectoryData = cumsum(draws[,1])/seq(1,1000)
34 plot(1:1000, TrajectoryData, type = "l", xlab = "Iterations", ylab = "Mean (u)", main = "
    Trajectories of the sampled Markov chains")
35
36 TrajectoryData = cumsum(draws[,2])/seq(1,1000)
37 plot(1:1000, TrajectoryData, type = "l", xlab = "Iterations", ylab = "Variance (sigma^2)", main = "
    Trajectories of the sampled Markov chains")
38
39 mean(draws[,1])
40 mean(draws[,2])
```

2.1.2 b

See NormalMixtureGibbs.R on course page. Changed nComp to 2.

2.1.3 c

```
1 u = 32.2129
2 sigma= 1546.929
3
4 plot(density(x), breaks = 20, freq = FALSE, xlim = c(xGridMin,xGridMax), main = "Fitted
   density")
5 lines(xGrid, mixDensMean, type = "l", lwd = 2, lty = 4, col = "red")
6 lines(xGrid, dnorm(xGrid, mean = u, sd = sqrt(sigma)), type = "l", lwd = 2, col = "blue")
7 legend("topright", box.lty = 1, legend = c("Data histogram","Mixture normal density","Normal
   density"), col=c("black","red","blue"), lwd = 2)
```

2.2 Assignment 2

```
1
2 set.seed(1234)
3 data <- read.table("eBayNumberOfBidderData.dat", header = T)
4
5 # a
6
7 model <- glm(nBids ~ . - Const, family="poisson", data=data)
8 barplot(model$coefficients, ylab = "Beta", xlab="Covariate", main="Covariates with their
   corresponding beta values")
9 summary(model)
10
11 # b
12
13 library(mvtnorm)
14 y <- data$nBids
15 X <- as.matrix(data[, -c(1)])
16
17 Sigma <- 100 * solve(t(X)%*%X)
18 mu <- matrix(0, nrow=1, ncol=9)
19
20
21 LogPoisson <- function(betas,y,X,mu,Sigma){
22   linPred <- betas %*% t(X)
23
24   logLik <- sum(linPred * y - exp(linPred))
25   if (abs(logLik) == Inf) logLik = -20000
26
27   logPrior <- dmvnorm(betas, matrix(0, length(betas), 1), Sigma, log=TRUE) ## zellners prior
28   return(logLik + logPrior)
29 }
30 initVal <- as.vector(rep(0,dim(X)[2]));
31
32 OptimResults<-optim(initVal,LogPoisson,gr=NULL,y,X,mu,Sigma,method=c("BFGS"),control=list(
   fnscale=-1),hessian=TRUE)
33
34 beta_tilde <- OptimResults$par
35 j_inv <- -solve(OptimResults$hessian)
36 beta_draws <- rmvnorm(10000,beta_tilde, j_inv )
```

```

37
38 par(mfrow=c(3,3))
39
40 hist(beta_draws[, 1], 100, xlab = "Intercept", main = "")
41 hist(beta_draws[, 2], 100, xlab = "PowerSeller", main = "")
42 hist(beta_draws[, 3], 100, xlab = "VerifyID", main = "")
43 hist(beta_draws[, 4], 100, xlab = "Sealed", main = "")
44 hist(beta_draws[, 5], 100, xlab = "Minblem", main = "")
45 hist(beta_draws[, 6], 100, xlab = "MajBlem", main = "")
46 hist(beta_draws[, 7], 100, xlab = "LargNeg", main = "")
47 hist(beta_draws[, 8], 100, xlab = "LogBook", main = "")
48 hist(beta_draws[, 9], 100, xlab = "MinBidShare", main = "")
49
50 par(mfrow=c(1,1))
51
52 # c
53
54 Metropolis <- function(c, N, summa, postFunction, ...) {
55   proposal_last <- as.vector(rep(0, dim(X)[2]))
56   thetas <- matrix(0, N, dim(X)[2])
57   acceptance <- 0
58   for (x in 1:N) {
59     proposal <- as.vector(rmvnorm(1, mean = proposal_last, sigma = c * summa))
60     p_theta <- postFunction(proposal, ...)
61     p_theta_last <- postFunction(proposal_last, ...)
62     alpha <- min(1, exp(p_theta - p_theta_last))
63     rand <- runif(1)
64     if (alpha > rand) {
65       proposal_last <- proposal
66       acceptance <- acceptance + 1
67     }
68     thetas[x, ] <- proposal_last
69   }
70   cat("Acceptance Ratio: ", acceptance / N)
71   return (thetas)
72 }
73
74 n <- 10000
75
76 draws <- Metropolis(0.5, n, j_inv, LogPoisson, y, X, mu, Sigma)
77
78 cum_mean <- matrix(0, n, dim(X)[2])
79
80 for (i in 1:dim(draws)[2]) {
81   cum_mean[,i] <- cumsum(draws[,i]) / seq_along(draws[,i])
82 }
83
84 par(mfrow=c(3,3))
85
86 plot(draws[,1], main="Intercept", xlab="iteration", ylab="Value", type="l")
87 plot(draws[,2], main="PowerSeller", xlab="iteration", ylab="Value", type="l")
88 plot(draws[,3], main="VerifyID", xlab="iteration", ylab="Value", type="l")
89 plot(draws[,4], main="Sealed", xlab="iteration", ylab="Value", type="l")
90 plot(draws[,5], main="Minblem", xlab="iteration", ylab="Value", type="l")
91 plot(draws[,6], main="MajBlem", xlab="iteration", ylab="Value", type="l")
92 plot(draws[,7], main="LargNeg", xlab="iteration", ylab="Value", type="l")
93 plot(draws[,8], main="LogBook", xlab="iteration", ylab="Value", type="l")
94 plot(draws[,9], main="MinBidShare", xlab="iteration", ylab="Value", type="l")
95
96 plot(cum_mean[,1], main="Intercept", xlab="iteration", ylab="Cumulative Mean", type="l")
97 abline(h=model$coefficients[1], col="red")
98 plot(cum_mean[,2], main="PowerSeller", xlab="iteration", ylab="Cumulative Mean", type="l")

```



```

99 abline(h=model$coefficients[2], col="red")
100 plot(cum_mean[,3], main="VerifyID", xlab="iteration", ylab="Cumulative Mean", type="l")
101 abline(h=model$coefficients[3], col="red")
102 plot(cum_mean[,4], main="Sealed", xlab="iteration", ylab="Cumulative Mean", type="l")
103 abline(h=model$coefficients[4], col="red")
104 plot(cum_mean[,5], main="Minblem", xlab="iteration", ylab="Cumulative Mean", type="l")
105 abline(h=model$coefficients[5], col="red")
106 plot(cum_mean[,6], main="MajBlem", xlab="iteration", ylab="Cumulative Mean", type="l")
107 abline(h=model$coefficients[6], col="red")
108 plot(cum_mean[,7], main="LargNeg", xlab="iteration", ylab="Cumulative Mean", type="l")
109 abline(h=model$coefficients[7], col="red")
110 plot(cum_mean[,8], main="LogBook", xlab="iteration", ylab="Cumulative Mean", type="l")
111 abline(h=model$coefficients[8], col="red")
112 plot(cum_mean[,9], main="MinBidShare", xlab="iteration", ylab="Cumulative Mean", type="l")
113 abline(h=model$coefficients[9], col="red")
114
115 par(mfrow=c(1,1))
116
117 betas_mean <- as.vector(draws[n-1,])
118 #d
119 cov = as.vector(c(1,1,1,1,0,0,0,1,0.5))
120
121 lambda = exp(cov%*%betas_mean)
122 nBids_sim = rpois(10000, lambda)
123
124 barplot(table(nBids_sim)/n, ylab="Probability", xlab="nBids")

```